

1. Prerequisites:	3
2. Installation of Kafka – 90 Mins	8
3. Installation Confluent Kafka (Local) – 30 Minutes	16
4. Basic Kafka Operations - CLI (Topic) – 30 Mins	17
5. Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins	24
6. Kafka Connector (File & JDBC) - 150 Minutes	27
7. Schema Registry - Manage Schemas for Topics – 30 Minutes	56
9. Errors	62
I. {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)	62
10. LOG verification + segment Sizing	64
11. Annexure Code:	65
II. DumplogSegment.....	65
III. Resources	66

Hardware:

8 GB RAM , 30 GB HDD , Centos 7 or above OS. Access to internet. (3 Machine each per participants)

Software Inventory:

- Zookeeper Version: apache-zookeeper-3.6.1-bin.tar

2 Kafka – Administration

- Apache kafka : 2.13-3.1.0
- jdk-8u45-linux-x64.tar.gz (Any jdk above 1.8)
- Eclipse for Linux. (Any Latest version for JEE Development)
- Status : Color is Verified

Last Updated: Feb 26 - 2022.

1. Prerequisites:

Option I

Start the VM using VM player and Logon to the server using telnet or directly in the VM console. Enter the root credentials to logon.

You can copy files from the host to VM using winscp.exe.

Option II.

Using docker:

Instantiate a container, kafkao.

You can copy files from the host to container using docker copy command.

Execute the following command, it will perform the following:

- Create a network : spark-net
- download the image, centos:7
- start a container with the name, kafkao and mount host folder in /opt

```
#docker network create --driver bridge spark-net
```

4 Kafka – Administration

```
#docker run --name kafkao --hostname kafkao -p 9092:9092 -p 8081:8081 -p 2181:2181 -p 9999:9999 -i -t --privileged --network spark-net -v /Users/henrypotsangbam/Documents/Docker:/opt centos:7 /usr/sbin/init
```

Optional:

```
#docker run --name kafkao --hostname kafkao -p 9092:9092 -p 9081:8081 -p 9082:8082 -p 3181:2181 -p 9990:9999 -i -t --privileged --network spark-net -v /Volumes/Samsung_T5/software/:/Software -v /Volumes/Samsung_T5/software/install/:/opt -v /Volumes/Samsung_T5/software/data/:/data centos:7 /usr/sbin/init
```

You can verify the container from a separate terminal:

```
(base) Henrys-MacBook-Air:~ henrypotsangbam$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
0945477d457b        ehenry0227/learning:spark-kafka-raw   "/bin/bash"         34 seconds ago    Up 33 seconds
ds                  0.0.0.0:6062->6062/tcp, 0.0.0.0:8081->8081/tcp
(base) Henrys-MacBook-Air:~ henrypotsangbam$
```

Perform the installation after this as a normal server.

Update the server.properties.

```
// Don't change the listeners address if the client also run in the host machine.
```

```
advertised.listeners=PLAINTEXT://localhost:9092
```

5 Kafka – Administration

After installation to start the zookeeper and kafka broker:

Open two separate terminals and execute the following

```
# /opt/kafka/bin/zookeeper-server-start.sh /opt/kafka/config/zookeeper.properties  
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

Kafka 2nd Cluster: (It will be used as DC2 for Kafka Mirroring)

```
#docker run --name kafka1 --hostname kafka1 -p 7092:9092 -p 7081:8081 -p 4181:2181 -p  
8999:9999 -i -t --privileged --network spark-net -v  
/Volumes/Samsung_T5/software/:/Software -v  
/Volumes/Samsung_T5/software/install/:/opt -v  
/Volumes/Samsung_T5/software/data/:/data centos:7 /usr/sbin/init
```

Install the kafka as done before if required or copy the kafka folder in different name to persevere the libraries.

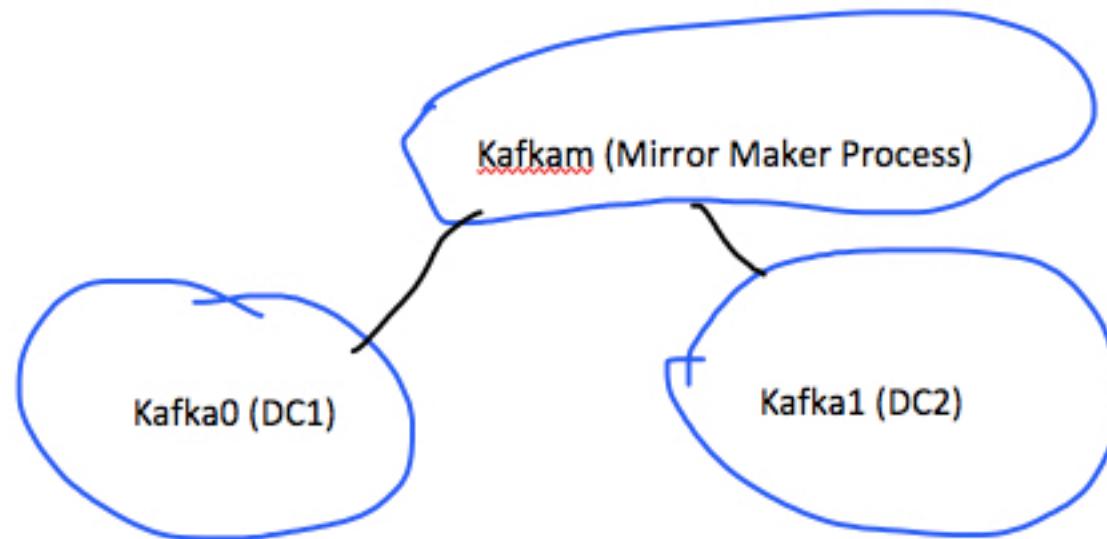
Kafka Mirror Maker: (It will execute the Kafka Mirroring process)

```
#docker run --name kafkam --hostname kafkam -p 6092:9092 -p 6081:8081 -i -t --  
privileged --network spark-net -v /Users/henrypotsangbam/Documents/Docker:/opt  
centos:7 /usr/sbin/init
```

Kafka Mirror Architecture:

6 Kafka – Administration

12:46 PM



Note:

If you are using docker ensure to update the server.properties with the following entries for accessing the broker from the host machine.

// Changes Begin

```
listeners=PLAINTEXT://0.0.0.0:9092,PLAINTEXT_HOST://0.0.0.0:8081  
advertised.listeners=PLAINTEXT://kafkao:9092,PLAINTEXT_HOST://localhost:8081  
listener.security.protocol.map=PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
```

// Changes End

use container, kafkao to connect from Docker. However, use localhost:8081 for connecting from the Host machine.

2. Installation of Kafka – 90 Mins

You need to install java before installing zookeeper and Kafka.

Installation of Java

```
#tar -xvf jdk-8u45-linux-x64.tar.gz -C /opt
```

Set the above path in the PATH variable and JAVA_HOME

```
export JAVA_HOME=/opt/jdk
```

```
export PATH=$PATH:$JAVA_HOME/bin
```

Include in the profile as follow:

```
#vi ~/.bashrc
```

9 Kafka – Administration

```
[root@tos opt]# more ~/.bashrc
# .bashrc

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

export JAVA_HOME=/opt/jdk1.8.0_45

export PATH=$PATH:$JAVA_HOME/bin
# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi
[root@tos opt]#
```

Update the shell scripts using the following command.

```
#bash
```

Installing Zookeeper

You can choose any of the options given below:

Option I (Fresh Installation): (We will use this for our lab)

The following steps install Zookeeper with a basic configuration in /opt/zookeeper.
Its configured to store data in /opt/data/zookeeper:

Extract the zookeeper archive file in /opt and rename the installation folder for brevity.

```
# tar -xvf apache-zookeeper-3.6.1-bin.tar -C /opt  
#mv a*zookeeper-3.6.1* /opt/zookeeper  
#mkdir -p /opt/data/zookeeper
```

Create a zookeeper configuration file and update with the following values.

```
#vi /opt/zookeeper/conf/zoo.cfg  
tickTime=2000  
dataDir=/opt/data/zookeeper  
clientPort=2181
```

Update the zoo.cfg with the above entries. You can save the file using esc+wq!

Start the zookeeper using the following scripts.

```
# /opt/zookeeper/bin/zkServer.sh start
```

```
[root@tos opt]# /opt/zookeeper/bin/zkServer.sh start  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/.../conf/zoo.cfg  
Starting zookeeper ... STARTED  
[root@tos opt]# █
```

Option II (Part of the Apache Kafka)

```
#bin/zookeeper-server-start.sh config/zookeeper.properties
```

You can now validate that Zookeeper is running correctly in standalone mode by connecting to the client port and sending the four-letter command `srvr`:

```
#telnet localhost 2181
```

```
➤ srvr
```

```
[root@tos opt]# telnet localhost 2181
Trying :1...
Connected to localhost.
Escape character is '^].
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/
27/2018 03:55 GMT
Latency min/avg/max: 0/0/0
Received: 2
Sent: 1
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 4
Connection closed by foreign host.
[root@tos opt]#
```

After zookeeper installation, let us install the Kafka Broker.

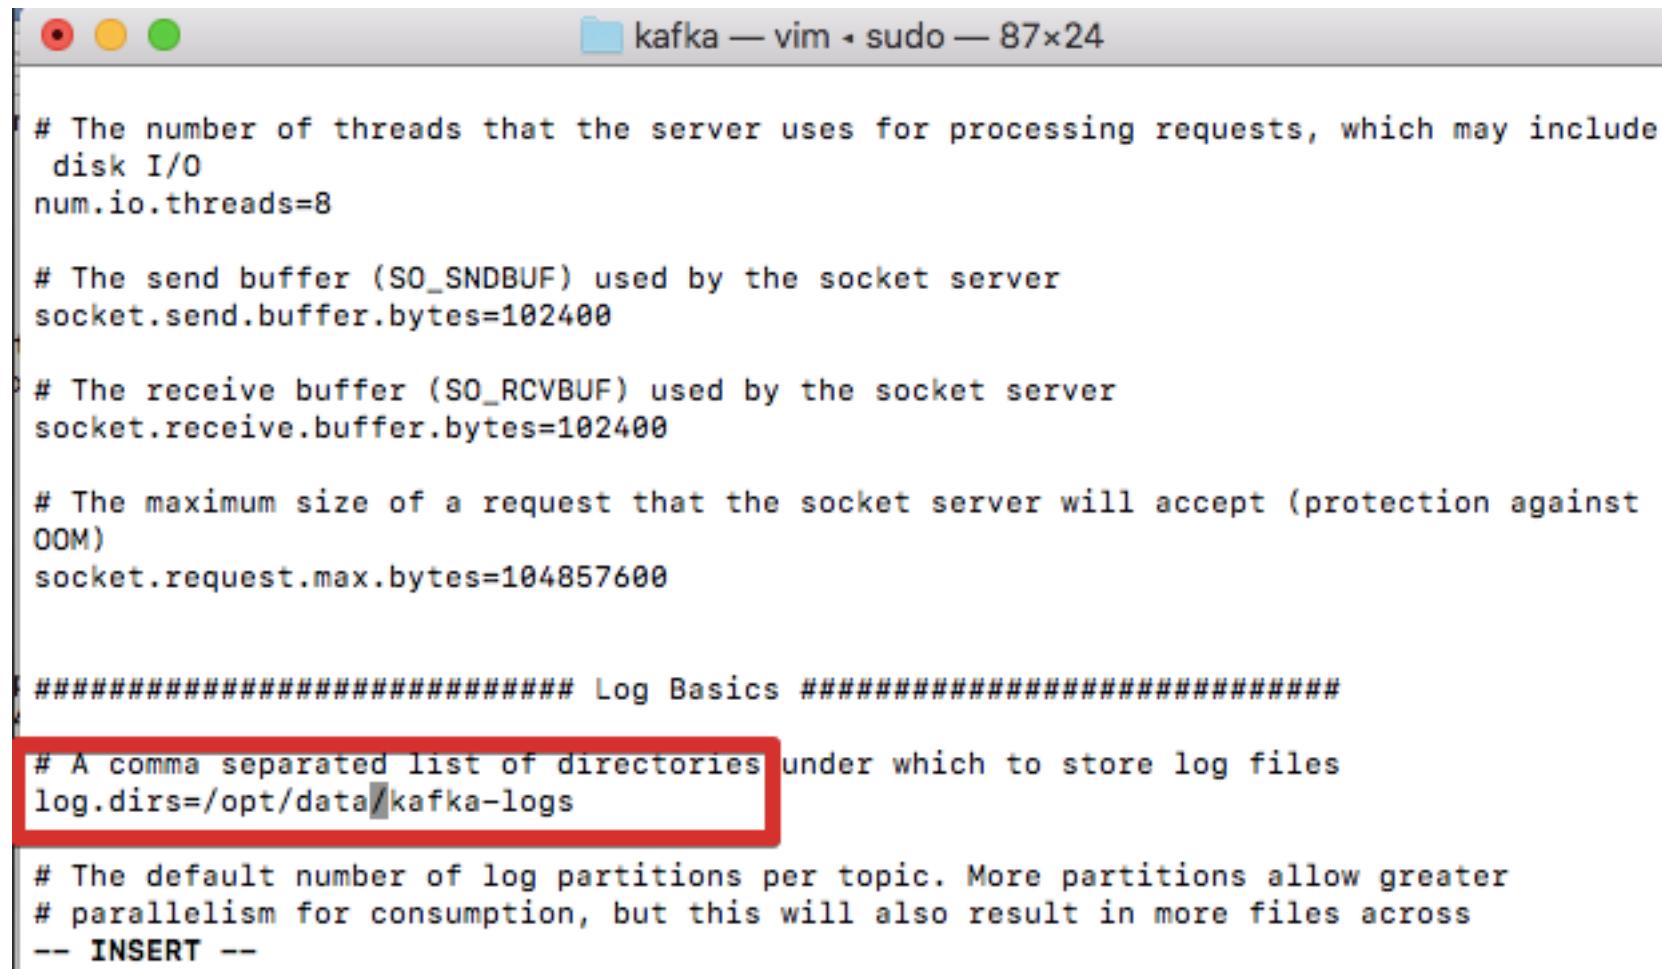
Installation of Kafka Broker

The following example installs Kafka in /opt/kafka, configured to use the

Zookeeper server started previously and to store the message log segments stored in /tmp/kafka-logs:

```
# tar -zxf kafka_2.12-* -C /opt
# mv kafka_2.1* /opt/kafka
# mkdir /opt/data/kafka-logs
```

Update the /opt/kafka/config/server.properties to store the Kafka Log in the above mention folder.



```
# The number of threads that the server uses for processing requests, which may include disk I/O
num.io.threads=8

# The send buffer (SO_SNDBUF) used by the socket server
socket.send.buffer.bytes=102400

# The receive buffer (SO_RCVBUF) used by the socket server
socket.receive.buffer.bytes=102400

# The maximum size of a request that the socket server will accept (protection against OOM)
socket.request.max.bytes=104857600

##### Log Basics #####
# A comma separated list of directories under which to store log files
log.dirs=/opt/data/kafka-logs

# The default number of log partitions per topic. More partitions allow greater
# parallelism for consumption, but this will also result in more files across
-- INSERT --
```

If you are using docker, kindly refer the prerequisite section for setting specific to docker.

Start the broker with the following command

```
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

```
[root@tos opt]# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
[root@tos opt]# jps
3476 Kafka
3499 Jps
2895 QuorumPeerMain
[root@tos opt]#
```

```
#mkdir /opt/scripts
```

All the common execution scripts will be stored in the above folder.

The following scripts will start a zookeeper along with a broker. Create the following files and update with the following scripts. It will start the zookeeper and kafka broker using the mention script.

```
#cd /opt/scripts
#vi startABroker.sh
##### Scripts Begin #####
#!/usr/bin/env bash

# Start Zookeeper.
/opt/zookeeper/bin/zkServer.sh start

#Start Kafka Server
/opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

```
echo "Started Successfully"
```

```
// ##### Scripts End  
#####
```

To shutdown the Broker, find the process and kill.

```
ps -eaf | grep java
```

or

create a scripts : /opt/scripts/stopBroker.sh with the following commands in it.

```
#!/usr/bin/env bash
```

```
/opt/kafka/bin/kafka-server-stop.sh
```

To Stop Zookeeper create the following script.

```
#vi /opt/scripts/stopZookeeper.sh
```

Update the following commands in the above script and save it.

```
#!/usr/bin/env bash
```

```
# Stop Zookeeper.
```

```
/opt/zookeeper/bin/zkServer.sh stop
```

```
echo "Stop zookeeper Successfully"
```

Lab Installation completes End here.

3. Installation Confluent Kafka (Local) – 30 Minutes

The purpose of this lab is to demonstrate the basic and most powerful capabilities of Confluent Platform – Schema Registry.

Installing Confluent kafka

Inflate the confluent kafka compress file as shown below:

```
#tar -xvf confluent-7.0.1.tar -C /opt
```

Set the environment variable for the Confluent Platform directory.

```
export CONFLUENT_HOME=/opt/confluent
```

Set your PATH variable:

```
# vi ~/.bashrc
```

```
export PATH=/opt/confluent/bin:${PATH};
```

----- Lab Ends Here -----

4. Basic Kafka Operations - CLI (Topic) – 30 Mins

In this lab you will be able to create a topic and perform some operations to understand the information about topic like partition and replication.

You need to start the broker using startABroker.sh. The script should be in /opt/scripts folder

```
#sh startABroker.sh  
#jps
```

```
[root@tos scripts]# sh startABroker.sh  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/..../conf/zoo.cfg  
Starting zookeeper ... STARTED  
Started Successfully  
[root@tos scripts]# jps  
11665 Jps  
11646 Kafka  
11375 QuorumPeerMain  
[root@tos scripts]#
```

Once the Kafka broker is started, we can verify that it is working by performing some simple operations against the broker; creating a test topic etc.

Create and verify details about topic:

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 1 --topic test
```

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 12 --topic IBM  
  
# /opt/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic test  
  
#/opt/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic IBM
```

```
[root@tos opt]# /opt/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic test  
Created topic "test".  
[root@tos opt]# /opt/kafka/bin/kafka-topics.sh --zookeeper localhost:2181 --describe --topic test  
Topic:test      PartitionCount:1      ReplicationFactor:1      Configs:  
          Topic: test      Partition: 0      Leader: 0      Replicas: 0      Isr: 0  
[root@tos opt]#
```

Verify the no of partition in the output.

list and describe topic.

```
#/opt/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --list
```

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --zookeeper localhost:2181 --list  
__consumer_offsets  
my-failsafe-topic  
test
```

List and describe Topics

What does the tool do?

This tool lists the information for a given list of topics. If no topics are provided in the command line, the tool queries zookeeper to get all the topics and lists the information for them. The fields that the tool displays are - topic name, partition, leader, replicas, isr.

How to use the tool?

List only single topic named "test" (prints only topic name)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092 --topic test
```

List all topics (prints only topic names)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

```
[root@tos scripts]# jps
12960 Jps
12314 Kafka
12043 QuorumPeerMain
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --list --zookeeper tos.master.com:2181 --topic
test
test
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --list --zookeeper localhost:2181
CustomerCountry
__consumer_offsets
henry-topic
my-failsafe-topic
my-kafka-topic
my-kafka-topic1
test
test-topic
```

Describe only single topic named "test" (prints details about the topic)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server localhost:9092 --topic test
```

Describe all topics (prints details about the topics)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server localhost:9092
```

```
[root@tos scripts]#  
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --describe --zookeeper localhost:2181 --topic test  
Topic:test PartitionCount:1 ReplicationFactor:1 Configs:  
  Topic: test Partition: 0 Leader: 0 Replicas: 0 Isr: 0  
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --describe --zookeeper localhost:2181  
Topic:CustomerCountry PartitionCount:1 ReplicationFactor:1 Configs:  
  Topic: CustomerCountry Partition: 0 Leader: 1 Replicas: 1 Isr: 1  
Topic:_consumer_offsets PartitionCount:50 ReplicationFactor:1 Configs:segment.bytes=104857600,cleanup.policy=compact,compression.type=producer  
  Topic: __consumer_offsets Partition: 0 Leader: 1 Replicas: 1 Isr: 1  
  Topic: __consumer_offsets Partition: 1 Leader: 2 Replicas: 2 Isr: 2
```

We will understand the output in details later.

Create Topics

What does the tool do?

By default, Kafka auto creates topic if "auto.create.topics.enable" is set to true on the server. This creates a topic with a default number of partitions, replication factor and uses Kafka's default scheme to do replica assignment. Sometimes, it may be required that we would like to customize a topic while creating it. This tool helps to create a topic and also specify the number of partitions, replication factor and replica assignment list for the topic.

How to use the tool?

create topic with default settings

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --topic topic1 --partitions 2 --replication-factor 1
```

```
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --topic topic1 --partitions 2 --replication-factor 1
Created topic "topic1".
[root@tos scripts]#
```

Create a topic with replication 2.

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --topic topic2 --partitions 2 --replication-factor 2
```

```
[root@kafka0 /]# /opt/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --topic topic2 --partitions 2 --replication-factor 2
Error while executing topic command : Replication factor: 2 larger than available brokers: 1.
[2022-01-10 07:44:51,716] ERROR org.apache.kafka.common.errors.InvalidReplicationFactorException: Replication factor: 2 larger than available brokers: 1.
(kafka.admin.TopicCommand$)
[root@kafka0 /]# /opt/kafka/bin/kafka-topics.sh --describe --zookeeper localhost:2181 --topic topic2
Error while executing topic command : Topic 'topic2' does not exist as expected
[2022-01-10 07:45:26,938] ERROR java.lang.IllegalArgumentException: Topic 'topic2' does not exist as expected
    at kafka.admin.TopicCommand$.kafka$admin$TopicCommand$$ensureTopicExists(TopicCommand.scala:542)
    at kafka.admin.TopicCommand$ZookeeperTopicService.describeTopic(TopicCommand.scala:447)
    at kafka.admin.TopicCommand$.main(TopicCommand.scala:69)
    at kafka.admin.TopicCommand.main(TopicCommand.scala)
(kafka.admin.TopicCommand$)
[root@kafka0 /]#
```

As shown above, it generates an error. Since there is only a single broker. It will fix later.

Lab CLI completes End here.

5. Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins

In this lab we will send message to the broker and consumer message using the kafka inbuilt commands.

You need to complete the previous lab before proceeding ahead.

You need to start the broker using startABroker.sh if not done earlier. The script should be in /opt/scripts folder

```
#sh startABroker.sh  
#jps
```

```
[root@tos scripts]# sh startABroker.sh  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/.../conf/zoo.cfg  
Starting zookeeper ... STARTED  
Started Successfully  
[root@tos scripts]# jps  
11665 Jps  
11646 Kafka  
11375 QuorumPeerMain  
[root@tos scripts]#
```

Sent message to **test** topic: Open a console to send message to the topic, test. Enter some text as shown below.

```
# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test  
Test Message 1  
Test Message 2  
^D
```

```
#
```

```
[root@tos config]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test  
>hi  
>Hello  
>TEst message  
>[root@tos config]#
```

Consume messages from a test topic: As soon as you enter the following script in a separate terminal, you should be able to consume the messages that we have type in the producer console.

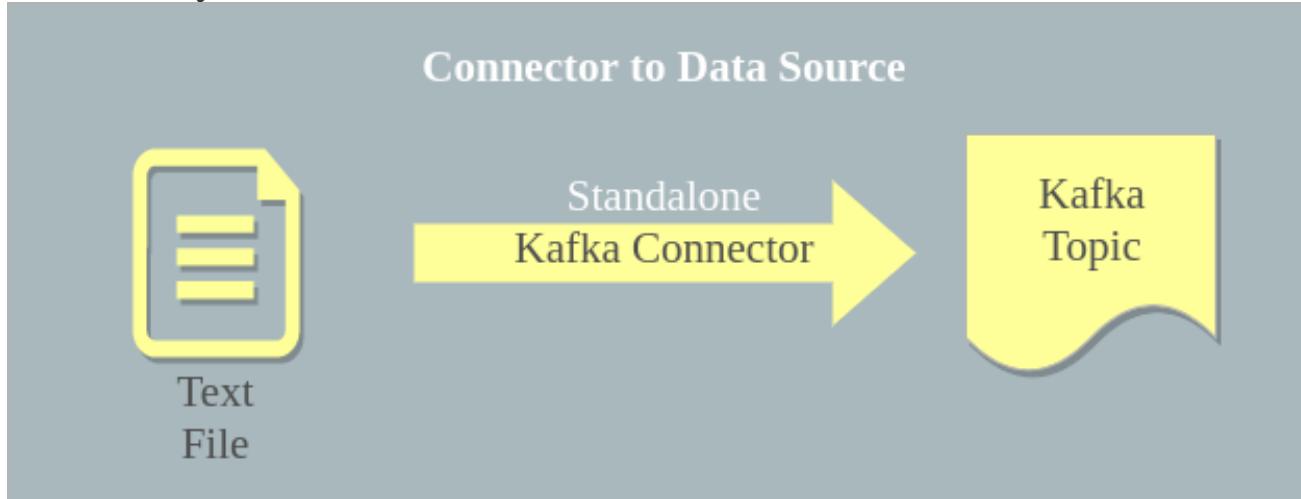
```
# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test  
--from-beginning
```

```
[base] Henrys-MacBook-Air:kafka henrypotsangbam$ /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test --from-beginning
Test Message 1
Test Message 2
Sending third message
```

Lab CLI ends here.

6. Kafka Connector (File & JDBC) - 150 Minutes

Apache Kafka Connector – Connectors are the components of Kafka that could be setup to listen the changes that happen to a data source like a file or database, and pull in those changes automatically.



When working with Kafka you might need to write data from a local file to a Kafka topic. This is actually very easy to do with Kafka Connect. Kafka Connect is a framework that provides scalable and reliable streaming of data to and from Apache Kafka. With Kafka Connect, writing a file's content to a topic requires only a few simple steps

Starting Kafka and Zookeeper

We will use the standalone Broker this example.

```
#cd /opt/scripts  
#sh startABroker.sh
```

```
[root@tos scripts]# ls
custom-reassignment.json  server2.log      stopZookeeper.sh
old.json                  start3Brokers.sh topics-to-move.json
server0.log                startABroker.sh  zookeeper.out
server1.log                stop3Brokers.sh
[root@tos scripts]# sh startABroker.sh
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/..../conf/zoo.cfg
Starting zookeeper ... STARTED
Started Successfully
[root@tos scripts]# jps
2209 QuorumPeerMain
2474 Kafka
2493 Jps
[root@tos scripts]#
```

Creating a Topic to Write to, the message will be fetch from the file and publish to the topic.

```
#cd /opt/scripts
#/opt/kafka/bin/kafka-topics.sh \
--create \
--bootstrap-server localhost:9092 \
--replication-factor 1 \
--partitions 1 \
--topic my-kafka-connect
```

```
[root@tos bin]#  
[root@tos bin]# cd /opt/scripts  
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh \  
> --create \  
> --zookeeper tos.master.com:2181 \  
> --replication-factor 1 \  
> --partitions 1 \  
> --topic my-kafka-connect  
Created topic "my-kafka-connect".  
[root@tos scripts]#
```

Creating a Source Config File

Since we are reading the contents of a local file and writing to Kafka, this file is considered our “source”. Therefore, we will use the FileSource connector. We must create a configuration file to use with this connector. For this most part you can copy the example available in `$KAFKA_HOME/config/connect-file-source.properties`. Below is an example of our `my-file-source.properties` file.

Open a new file.

```
#vi my-file-source.properties
```

Paste the following instruction in the above file.

```
#my-file-source.properties config file
name=local-file-source
connector.class=FileStreamSource
tasks.max=1
file=/tmp/my-test.txt
topic=my-kafka-connect
```

This file indicates that we will use the FileStreamSource connector class, read data from the /tmp/my-test.txt file, and publish records to the my-kafka-connect Kafka topic. We are also only using 1 task to push this data to Kafka, since we are reading/publishing a single file.

Creating a Worker Config File.

Processes that execute Kafka Connect connectors and tasks are called workers. Since we are reading data from a single machine and publishing to Kafka, we can use the simpler of the two types, standalone workers (as opposed to distributed workers). You can find a sample config file for standalone workers in \$KAFKA_HOME/config/connect-standalone.properties. We will call our file my-standalone.properties.

Create a file.

```
#vi my-standalone.properties
```

Update with the following contents.

```
#bootstrap kafka servers  
bootstrap.servers=localhost:9092
```

```
# specify input data format  
key.converter=org.apache.kafka.connect.storage.StringConverter  
value.converter=org.apache.kafka.connect.storage.StringConverter
```

```
# The internal converter used for offsets, most will always want to use the built-in default  
internal.key.converter=org.apache.kafka.connect.json.JsonConverter  
internal.value.converter=org.apache.kafka.connect.json.JsonConverter  
internal.key.converter.schemas.enable=false  
internal.value.converter.schemas.enable=false
```

```
# local file storing offsets and config data  
offset.storage.file.filename=/tmp/connect.offsets
```

The main change in this example in comparison to the default is the key.converter and value.converter settings. Since our file contains simple text, we use the StringConverter types.

Running Kafka Connect.

Now it is time to run Kafka Connect with our worker and source configuration files. As mentioned before we will be running Kafka Connect in standalone mode. Here is an example of doing this with our custom config files:

```
#/opt/kafka/bin/connect-standalone.sh my-standalone.properties my-file-source.properties
```

```
mSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:35,218] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:36,226] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:37,241] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:38,250] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:39,256] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:40,258] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)  
[2018-06-17 12:55:41,272] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)
```

Open a terminal.

Our input file /tmp/my-test.txt will be read in a single process to the Kafka my-kafka-connect topic. Here is a look at the file contents:

Create a file and paste the following text.

```
#vi /tmp/my-test.txt
```

This Message is from Test File.

It will be consumed by the Kafka Connector.

```
[2018-06-17 12:56:53,991] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)
[2018-06-17 12:56:55,535] INFO Cluster ID: rwk6R3n9TYSCArCsSs8V4w (org.apache.kafka.clients.Metadata:265)
[2018-06-17 12:57:53,852] INFO WorkerSourceTask{id=local-file-source-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:328)
[2018-06-17 12:57:53,854] INFO WorkerSourceTask{id=local-file-source-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:345)
[2018-06-17 12:57:54,062] INFO WorkerSourceTask{id=local-file-source-0} Finished commitOffsets successfully in 208 ms (org.apache.kafka.connect.runtime.WorkerSourceTask:427)
```

Open another terminal and execute the following consumer to consume the message.

Reading from the Kafka Topic

If we read from the Kafka topic that we created earlier, we should see the 2 lines in the source file that were written to Kafka:

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic my-kafka-connect --from-beginning
```

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server tos.master.com:9092 --topic my-kafka-connect --from-beginning
This Message is from Test File.
It will be consumed by the Kafka Connector.
```

Update the /tmp/my-test.txt with a new line “Hope this message arrived on the console.” . It should be shown on the consumer console automatically.

The screenshot shows two terminal windows. The top window displays the contents of a file named 'my-test.txt' which contains a test message for a Kafka connector. The bottom window shows the command used to start the Kafka connector.

```
login as: root
root@10.10.20.24's password:
Last login: Sun Jun 17 12:28:03 2018 from 10.10.20.1
[root@tos ~]# vi /tmp/my-test.txt
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server tos.master.com:9092 --topic my-kafka-connect --from-beginning
This Message is from Test File.
It will be consumed by the Kafka Connector.
Hope this message arrived on the console.
root@tos:~
```

```
root@tos:~
```

```
login as: root
root@10.10.20.24's password:
Last login: Sun Jun 17 12:56:11 2018 from 10.10.20.1
[root@tos ~]# vi /tmp/my-test.txt
[root@tos ~]#
```

Let us try with json file.

Stop the kafka connector – ctrl + C of the terminal in which we have started it.

Update the source properties files as shown below: (/opt/scripts/my-file-source.properties)

The screenshot shows a terminal window displaying the contents of a file named 'my-file-source.properties'. The file contains configuration settings for a Kafka connector, specifically for reading from a local file source.

```
#my-file-source.properties config file
name=local-file-source
connector.class=FileStreamSource
tasks.max=1
file=/tmp/my-test.json
topic=my-kafka-connect
~
```

Start the Kafka connector now:

```
#/opt/kafka/bin/connect-standalone.sh my-standalone.properties my-file-source.properties
```

Update the /tmp/my-test.json with the following:

```
{  
  "name": "John",  
  "age": 30,  
  "cars": [ "Ford", "BMW", "Fiat" ]  
}
```

You will be able to see the above json document in the consumer console. Start the consumer if you have closed it earlier.

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server tos.master.com:9092 --topic my-kafka-connect --from-beginning  
This Message is from Test File.  
It will be consumed by the Kafka Connector.  
Hope this message arrived on the console.  
{  
  "name": "John",  
  "age": 30,  
  "cars": [ "Ford", "BMW", "Fiat" ]  
},
```

We have successfully configured, file connector in standalone mode.

Next, let us configure JDBC in a distributed mode.

Kafka Connect – JDBC Connector.

Download the connector from the following url.

Tools - <https://www.confluent.io/hub/confluentinc/kafka-connect-jdbc>

Download installation

Or download the ZIP file and extract it into one of the directories that is listed on the Connect worker's plugin.path configuration properties. This must be done on each of the installations where Connect will be run.

[Download](#)

Extract the file:

```
# yum install unzip  
# unzip confluent*zip
```

```
[root@4f5607f478bd software]# pwd  
/software  
[root@4f5607f478bd software]# ls  
README.md confluentinc-kafka-connect-jdbc-10.0.0 hbase-2.3.2-bin.tar log.txt master.txt  
[root@4f5607f478bd software]# ls confluentinc-kafka-connect-jdbc-10.0.0/  
assets doc etc lib manifest.json  
[root@4f5607f478bd software]#
```

Rename the folder:

```
#mv confluentinc-kafka-connect-jdbc-10.0.0/ /software/plugins/kafka-jdbc
```

Add this to the plugin path in your Connect properties file.

For example, `plugin.path=/software/plugins`. Kafka Connect finds the plugins using its plugin path.

Start the Connect workers with that configuration. Connect will discover all connectors defined within those plugins.

We can create create /software/connect-distributed.properties file to specify the worker properties as follows:

Note that the `plugin.path` is the path that we need to place the library that we downloaded.

```
#vi /software/connect-distributed.properties
```

Update the following content in the above file.

```
# A list of host/port pairs to use for establishing the initial connection to the Kafka
cluster.
bootstrap.servers=localhost:9092

# unique name for the cluster, used in forming the Connect cluster group. Note that this
must not conflict with consumer group IDs
group.id=connect-cluster

# The converters specify the format of data in Kafka and how to translate it into Connect
data.
key.converter=org.apache.kafka.connect.json.JsonConverter
value.converter=org.apache.kafka.connect.json.JsonConverter
key.converter.schemas.enable=true
value.converter.schemas.enable=true

# Topic to use for storing offsets. This topic should have many partitions and be replicated
and compacted.
offset.storage.topic=connect-offsets
offset.storage.replication.factor=1

# Topic to use for storing connector and task configurations; note that this should be a
single partition, highly replicated,
config.storage.topic=connect-configs
config.storage.replication.factor=1

# Topic to use for storing statuses. This topic can have multiple partitions and should be
replicated and compacted.
status.storage.topic=connect-status
status.storage.replication.factor=1

# Flush much faster than normal, which is useful for testing/debugging
offset.flush.interval.ms=10000

plugin.path=/software/plugins
```

Note that the `plugin.path` is the path that we need to place the library that we downloaded.

Start the kafka connector

Go to the bin folder of the kafka installation or specify the full path.

```
#/opt/kafka/bin/connect-distributed.sh /software/connect-distributed.properties
```

```
nnectorIds=[], taskIds=[], revokedConnectorIds=[], revokedTaskIds=[], delay=0} with rebalance delay: 0 (org.apache.kafka.connect.runtime.distr
ibuted.DistributedHerder:1681)
[2020-11-15 12:59:43,196] INFO [Worker clientId=connect-1, groupId=connect-cluster] Starting connectors and tasks using config offset -1 (org.
apache.kafka.connect.runtime.distributed.DistributedHerder:1208)
[2020-11-15 12:59:43,196] INFO [Worker clientId=connect-1, groupId=connect-cluster] Finished starting connectors and tasks (org.apache.kafka.c
onnect.runtime.distributed.DistributedHerder:1236)
Nov 15, 2020 12:59:43 PM org.glassfish.jersey.internal.Errors logErrors
WARNING: The following warnings have been detected: WARNING: The (sub)resource method listLoggers in org.apache.kafka.connect.runtime.rest.res
ources.LoggingResource contains empty path annotation.
WARNING: The (sub)resource method createConnector in org.apache.kafka.connect.runtime.rest.resources.ConnectorsResource contains empty path an
notation.
WARNING: The (sub)resource method listConnectors in org.apache.kafka.connect.runtime.rest.resources.ConnectorsResource contains empty path ann
otation.
WARNING: The (sub)resource method listConnectorPlugins in org.apache.kafka.connect.runtime.rest.resources.ConnectorPluginsResource contains em
pty path annotation.
WARNING: The (sub)resource method serverInfo in org.apache.kafka.connect.runtime.rest.resources.RootResource contains empty path annotation.

[2020-11-15 12:59:43,307] INFO Started o.e.j.s.ServletContextHandler@c6da8bb{/null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandle
r:825)
[2020-11-15 12:59:43,307] INFO REST resources initialized; server is started and ready to handle requests (org.apache.kafka.connect.runtime.re
st.RestServer:319)
[2020-11-15 12:59:43,308] INFO Kafka Connect started (org.apache.kafka.connect.runtime.Connect:57)
[2020-11-15 12:59:43,419] INFO [Worker clientId=connect-1, groupId=connect-cluster] Session key updated (org.apache.kafka.connect.runtime.dist
ributed.DistributedHerder:1570)
```

After running the connector we can confirm that connector's REST endpoint is accessible, and we can confirm that JDBC connector is in the plugin list by calling <http://localhost:8083/connector-plugins>

```
# curl http://localhost:8083/connector-plugins
```

```
[root@4f5607f478bd software]# curl http://localhost:8083/connector-plugins
[{"class":"io.confluent.connect.jdbc.JdbcSinkConnector","type":"sink","version":"10.0.0"}, {"class":"io.confluent.connect.jdbc.JdbcSourceConnector","type":"source","version":"10.0.0"}, {"class":"org.apache.kafka.connect.file.FileStreamSinkConnector","type":"sink","version":"2.6.0"}, {"class":"org.apache.kafka.connect.file.FileStreamSourceConnector","type":"source","version":"2.6.0"}, {"class":"org.apache.kafka.connect.mirror.MirrorCheckpointConnector","type":"source","version":"1"}, {"class":"org.apache.kafka.connect.mirror.MirrorHeartbeatConnector","type":"source","version":"1"}, {"class":"org.apache.kafka.connect.mirror.MirrorSourceConnector","type":"source","version":"1"}][root@4f5607f478bd software]#
[root@4f5607f478bd software]#
```

Install the postgress db. We will configure postgressdb to store some data which will be transfer by connect to kafka topic.

Refer url for downloading the postgresdb software : <https://www.postgresql.org/download/>.

For Centos 7:

[If you are using docker, initiate a container for PostgressDB:

```
#docker run -it --name postgresdb --privileged -p 5432:5432 centos:7 /usr/sbin/init
# docker exec -it postgresdb bash
```

Add both the container in same network.

```
#docker network connect spark-net postgresdb
]
```

The following command will install and initialize the postgres DB

```
#yum install -y https://download.postgresql.org/pub/repos/yum/reporpms/EL-7-x86_64/pgdg-redhat-repo-latest.noarch.rpm
#yum install -y postgresql13-server
#/usr/pgsql-13/bin/postgresql-13-setup initdb
```

Enable and start the service:

```
systemctl enable postgresql-13  
systemctl start postgresql-13
```

```
Complete!  
[root@5a1225a251e4 /]# /usr/pgsql-13/bin/postgresql-13-setup initdb  
Initializing database ... OK  
  
[root@5a1225a251e4 /]# systemctl enable postgresql-13  
Created symlink from /etc/systemd/system/multi-user.target.wants/postgresql-13.service to /usr/lib/systemd/system/postgresql-13.service.  
[root@5a1225a251e4 /]# systemctl start postgresql-13  
[root@5a1225a251e4 /]# systemctl status postgresql-13  
● postgresql-13.service - PostgreSQL 13 database server  
  Loaded: loaded (/usr/lib/systemd/system/postgresql-13.service; enabled; vendor preset: disabled)  
  Active: active (running) since Sun 2020-11-15 13:21:23 UTC; 6s ago  
    Docs: https://www.postgresql.org/docs/13/static/  
  Process: 322 ExecStartPre=/usr/pgsql-13/bin/postgresql-13-check-db-dir ${PGDATA} (code=exited, status=0/SUCCESS)  
 Main PID: 327 (postmaster)  
   CGroup: /docker/5a1225a251e4011bfe46ec657ce6ede8168c9c54e77cecc29df0c095aeebbd37/docker/5a1225a251e4011bfe46ec657ce6ede8168c9c54e77cecc29df0c095aeebbd37/system.slice/postgresql-13.service  
         ├ 327 /usr/pgsql-13/bin/postmaster -D /var/lib/pgsql/13/data/  
         ├ 328 postgres: logger  
         ├ 330 postgres: checkpointer  
         └ 331 postgres: background writer
```

The configuration for the plugin is stored in /software/jdbc-source.json file on kafka broker server. It contents is as follows: (replace localhost with container IP of postgress DB if you are using docker)

```
{
  "name": "jdbc_source_connector_postgresql_01",
  "config": {
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",
    "connection.url": "jdbc:postgresql://localhost:5432/postgres",
    "connection.user": "postgres",
    "connection.password": "postgres",
    "topic.prefix": "postgres-01-",
    "poll.interval.ms": 3600000,
    "mode": "bulk"
  }
}
```

Start Postgress DB CLI: On the postgres DB server.

```
sudo -u postgres psql
ALTER USER postgres PASSWORD 'postgres';
```

```
complete.
[root@5a1225a251e4 /]# sudo -u postgres psql
psql (13.1)
Type "help" for help.

postgres=# ALTER USER postgres PASSWORD 'postgres';
ALTER ROLE
postgres#
```

if sudo doesn't exists in your system, install (yum install sudo)

Create a table and insert few records: Execute all the bellows SQL commands from the Postgresdb CLI.

```
#CREATE TABLE leads (id INTEGER PRIMARY KEY, name VARCHAR);
insert into leads values (1,'Henry P');
insert into leads values (2,'Rajnita P');
insert into leads values (3,'Henderson P');
insert into leads values (4,'Tiraj P');
select * from leads;
```

```
postgres=# insert into leads values (4,'Tiraj P');
INSERT 0 1
postgres=# select * from leads;
 id |      name
----+
  1 | Henry P
  2 | Rajnita P
  3 | Henderson P
  4 | Tiraj P
(4 rows)
```

To exit : \q

Allow all to connect to postgresDB from remote server by updating the following contents:

```
#vi /var/lib/pgsql/13/data/postgresql.conf
```

```
listen_addresses = '*'
```

Append in the last line of the following file.

```
# vi /var/lib/pgsql/13/data/pg_hba.conf
```

```
host all all all trust
```

```
#systemctl restart postgresql-13
```

Starting the JDBC Connector – On the Kafka Node.

As we operate on distributed mode we run the connectors by calling REST endpoints with the configuration JSON. We can specify the configuration payload from a file for curl command. The following command starts the connector. Execute the following command from the directory you have stored the configuration json file.

```
curl -d @"jdbc-source.json" \  
-H "Content-Type: application/json" \  
-X POST http://localhost:8083/connectors
```

```
timestamp.column.name = □
timestamp.delay.interval.ms = 0
timestamp.initial = null
topic.prefix = postgres-02-
validate.non.null = true
(io.confluent.connect.jdbc.source.JdbcSourceTaskConfig:354)
[2020-11-15 16:05:29,313] INFO Using JDBC dialect PostgreSQL (io.confluent.connect.jdbc.source.JdbcSourceTask:98)
[2020-11-15 16:05:29,317] INFO Attempting to open connection #1 to PostgreSQL (io.confluent.connect.jdbc.util.CachedConnectionProvider:82)
[2020-11-15 16:05:29,383] INFO Started JDBC source task (io.confluent.connect.jdbc.source.JdbcSourceTask:257)
[2020-11-15 16:05:29,384] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Source task finished initialization and start (org.apache.kafka.connect.runtime.WorkerSourceTask:233)
[2020-11-15 16:05:29,388] INFO Begin using SQL query: SELECT * FROM "public"."leads" (io.confluent.connect.jdbc.source.TableQuerier:164)
[2020-11-15 16:05:29,503] WARN [Producer clientId=connector-producer-jdbc_source_connector_postgresql_04-0] Error while fetching metadata with correlation id 3 : {postgres-02-leads=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient:1073)
[2020-11-15 16:05:39,266] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:478)
[2020-11-15 16:05:39,268] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:495)
[2020-11-15 16:05:39,295] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Finished commitOffsets successfully in 28 ms (org.apache.kafka.connect.runtime.WorkerSourceTask:574)
[2020-11-15 16:05:49,264] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:478)
[2020-11-15 16:05:49,265] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:495)
```

We can see that in postgres database table leads is loaded to kafka topic- postgres-02-leads:

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

```
[root@4f5607f478bd software]# /opt/kafka/bin/kafka-topics.sh --list --zookeeper localhost:2181
__consumer_offsets
connect-configs
connect-offsets
connect-status
postgres-02-leads
test
[root@4f5607f478bd software]#
```

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server localhost:9092 --topic
postgres-01-leads
```

```
Topic: connect-status    Partition: 4    Leader: 0      Replicas: 0      Isr: 0
Topic: postgres-02-leads    PartitionCount: 1      ReplicationFactor: 1      Configs:
Topic: postgres-02-leads    Partition: 0    Leader: 0      Replicas: 0      Isr: 0
Topic: test    PartitionCount: 1      ReplicationFactor: 1      Configs:
Topic: test    Partition: 0    Leader: 0      Replicas: 0      Isr: 0
[root@4f5607f478bd software]#
```

And each row in the tables are loaded as a message. You can verify using the consumer console.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-01-leads --from-beginning
```

```
[root@4f5607f478bd software]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-02-leads --from-beginning
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}], "optional": false, "name": "leads"}, "payload": {"id": 1, "name": "Henry P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}], "optional": false, "name": "leads"}, "payload": {"id": 2, "name": "Rajnita P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}], "optional": false, "name": "leads"}, "payload": {"id": 3, "name": "Henderson P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}], "optional": false, "name": "leads"}, "payload": {"id": 4, "name": "Tiraj P"}}
```

Let us perform a Single Message Transformation.

Scenario: **Name** field should be changed to **FullName**

Solution: We will use the ReplaceField to perform the above activity.

Add another instance of connector.

Define the Single Message Transformation parameters.

The configuration for the plugin is stored in /software/jdbc-source_smt.json file on kafka broker server. It contents is as follows: (replace localhost with container IP of postgress DB if you are using docker)

```
{  
    "name": "jdbc_source_connector_postgresql_02",  
    "config": {  
        "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",  
        "connection.url": "jdbc:postgresql://localhost:5432/postgres",  
        "connection.user": "postgres",  
        "connection.password": "postgres",  
        "topic.prefix": "postgres-02-",  
        "poll.interval.ms" : 3600000,  
        "mode": "bulk",  
        "transforms": "RenameField",  
        "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value",  
        "transforms.RenameField.renames": "name:fullname"  
    }  
}
```

In the above configuration, the field , name is replace with fullname attribute.

Start the instance

```
#curl -d @"jdbc-source_smt1.json" \  
-H "Content-Type: application/json" \  
-X POST http://localhost:8083/connectors
```

Verify the topic and check the message inside that topic.

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-02-leads --from-beginning
```

```
[root@kafka0 software]# /opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
__consumer_offsets  
connect-configs  
connect-offsets  
connect-status  
postgres-01-leads  
postgres-02-leads  
[root@kafka0 software]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-02-leads --from-beginning  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}]}, "optional": false, "name": "leads"}, "payload": {"id": 1, "fullname": "Henry P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}]}, "optional": false, "name": "leads"}, "payload": {"id": 2, "fullname": "Rajnita P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}]}, "optional": false, "name": "leads"}, "payload": {"id": 3, "fullname": "Henderson P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}]}, "optional": false, "name": "leads"}, "payload": {"id": 4, "fullname": "Tiraj P"}}
```

As shown above the message field has been transform from Name to fullname.

Execute some Queries to get information about connectors:

Get a list of active connectors.

```
#curl http://localhost:8083/connectors
```

Get configuration info for a connector.

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/config
```

```
[root@kafka0 /]# curl http://localhost:8083/connectors
[{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "topic.prefix": "postgres-01-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "name": "jdbc_source_connector_postgresql_01", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}][root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/config
{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "topic.prefix": "postgres-01-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "name": "jdbc_source_connector_postgresql_01", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}[root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/config
{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "transforms.RenameField.renames": "name:fullname", "topic.prefix": "postgres-02-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value", "transforms": "RenameField", "name": "jdbc_source_connector_postgresql_02", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}[root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/tasks
[{"id": {"connector": "jdbc_source_connector_postgresql_02", "task": 0}, "config": {"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "transforms.RenameField.renames": "name:fullname", "connection.password": "postgres", "transforms": "RenameField", "mode": "bulk", "topic.prefix": "postgres-02-", "tables": "\\"public\\\".\\"leads\\\"", "task.class": "io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user": "postgres", "poll.interval.ms": "3600000", "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value", "name": "jdbc_source_connector_postgresql_02", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}][root@kafka0 /]#
```

Retrieve details for specific tasks (JDBC Tasks) - Get a list of tasks currently running for the connector.

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_02/tasks
```

```
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/tasks
[{"id": {"connector": "jdbc_source_connector_postgresql_02", "task": 0}, "config": {"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "transforms.RenameField.renames": "name:fullname", "connection.password": "postgres", "transforms": "Rename Field", "mode": "bulk", "topic.prefix": "postgres-02-", "tables": "\"public\".\"leads\"", "task.class": "io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user": "postgres", "poll.interval.ms": "3600000", "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value", "name": "jdbc_source_connector_postgresql_02", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}]}][root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/tasks
[{"id": {"connector": "jdbc_source_connector_postgresql_01", "task": 0}, "config": {"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "topic.prefix": "postgres-01-", "tables": "\"public\".\"leads\"", "connection.password": "postgres", "task.class": "io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user": "postgres", "poll.interval.ms": "3600000", "name": "jdbc_source_connector_postgresql_01", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}}][root@kafka0 /]#
```

The current status of the connector

```
#curl http://localhost:8083/connectors/?expand=status
```

```
[root@kafka0 /]# curl http://localhost:8083/connectors/?expand=status
{"jdbc_source_connector_postgresql_01": {"status": {"name": "jdbc_source_connector_postgresql_01", "connector": {"state": "RUNNING", "worker_id": "172.18.0.2:8083"}, "tasks": [{"id": 0, "state": "RUNNING", "worker_id": "172.18.0.2:8083"}], "type": "source"}}, "jdbc_source_connector_postgresql_02": {"status": {"name": "jdbc_source_connector_postgresql_02", "connector": {"state": "RUNNING", "worker_id": "172.18.0.2:8083"}, "tasks": [{"id": 0, "state": "RUNNING", "worker_id": "172.18.0.2:8083"}]}, "type": "source"}]}[root@kafka0 /]#
```

Gets the current status of the connector, including:

- Whether it is running or restarting, or if it has failed or paused
- Which worker it is assigned to
- Error information if it has failed
- The state of all its tasks

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/status
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/status
{"name":"jdbc_source_connector_postgresql_01","connector":{"state":"RUNNING","worker_id":"172.18.0.2:8083"},"tasks":[{"id":0,"state":"RUNNING","worker_id":"172.18.0.2:8083"}],"type":"source"}[root@kafka0 ~]#
```

Retrieve details for specific tasks

```
# curl
http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/tasks/0/status
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/tasks/0/status
{"id":0,"state":"RUNNING","worker_id":"172.18.0.2:8083"}[root@kafka0 ~]#
```

----- Labs End Here -----

7. Schema Registry - Manage Schemas for Topics – 30 Minutes

Pre-requisite: Install kafka server.

Download and install confluent kafka : Schema Registry only.

Start the Zookeeper and Broker.

Update the following in /opt/confluent/etc/schema-registry/schema-registry.properties

kafkastore.bootstrap.servers=PLAINTEXT://kafka0:9092

Start the Registry:

```
#schema-registry-start /opt/confluent/etc/schema-registry/schema-registry.properties
```

```
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.ModeResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.ModeResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SubjectsResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SubjectsResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SubjectVersionsResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SubjectVersionsResource will be ignored.  
Sep 28, 2020 7:47:55 AM org.glassfish.jersey.internal.inject.Providers checkProviderRuntime  
WARNING: A provider io.confluent.kafka.schemaregistry.rest.resources.SchemasResource registered in SERVER runtime does not implement any provider interfaces applicable in the SERVER runtime. Due to constraint configuration problems the provider io.confluent.kafka.schemaregistry.rest.resources.SchemasResource will be ignored.  
[2020-09-28 07:47:56,447] INFO HV000001: Hibernate Validator 6.0.17.Final (org.hibernate.validator.internal.util.Version:21)  
[2020-09-28 07:47:57,161] INFO Started o.e.j.s.ServletContextHandler@6b3871d6{/null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandler:825)  
[2020-09-28 07:47:57,215] INFO Started o.e.j.s.ServletContextHandler@aa10649{/ws>null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandler:825)  
[2020-09-28 07:47:57,272] INFO Started NetworkTrafficServerConnector@186f8716{HTTP/1.1,[http/1.1]}{0.0.0.0:8081} (org.eclipse.jetty.server.AbstractConnector:330)  
[2020-09-28 07:47:57,273] INFO Started @11223ms (org.eclipse.jetty.server.Server:399)  
[2020-09-28 07:47:57,275] INFO Server started, listening for requests... (io.confluent.kafka.schemaregistry.rest.SchemaRegistryMain:44)
```

<https://docs.confluent.io/platform/current/schema-registry/develop/using.html>

Create a topic and let us bind schema to this.

```
# kafka-topics.sh --create --bootstrap-server kafka0:9092 --partitions 1 --replication-factor 1 --topic my-kafka  
#kafka-topics.sh --list --bootstrap-server kafka0:9092
```

Use the Schema Registry API to add a schema for the topic my-kafka.

```
#curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data '{"schema": {"\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {\"name\": \"amount\", \"type\": \"double\"}]} }' http://localhost:8081/subjects/my-kafka-value/versions
```

```
[root@kafka0 code]# kafka-topics.sh --list --bootstrap-server kafka0:9092
__consumer_offsets
_schemas
my-kafka
[root@kafka0 code]# curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" --data '{"schema": {"\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {\"name\": \"amount\", \"type\": \"double\"}]} }' http://localhost:8081/subjects/my-kafka-value/versions
{"id":1}[root@kafka0 code]#
```

```
#curl -X GET http://localhost:8081/schemas/ids/1
```

```
[root@kafka0 code]# curl -X GET http://localhost:8081/schemas/ids/1
{"schema": {"\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {\"name\": \"amount\", \"type\": \"double\"}]} }[root@kafka0 code]#
```

You can verify the topic which maintains the schema details in the broker.

```
# kafka-topics.sh --list --bootstrap-server kafka0:9092
```

```
[root@kafka0 ~]# kafka-topics.sh --list --bootstrap-server kafka0:9092
__consumer_offsets
__schemas
my-kafka
transactions
[root@kafka0 ~]#
```

List all subjects associated with a given ID

Use this two-step process to find subjects associated with a given ID:

List all the subjects.

```
#curl -X GET http://localhost:8081/subjects
```

Iterate through the output from the subject list as follows, and check for the ID in the results:

```
#curl -X GET http://localhost:8081/subjects/<INSERT SUBJECT NAME>/versions/latest
```

E.x

```
#curl -X GET http://localhost:8081/subjects/my-kafka-value/versions/latest
```

```
[root@kafka0 code]# curl -X GET http://localhost:8081/subjects  
["my-kafka-value"] [root@kafka0 code]#  
[root@kafka0 code]# curl -X GET http://localhost:8081/subjects/my-kafka-value/versions/latest  
{"subject": "my-kafka-value", "version": 2, "id": 1, "schema": "{\"type\": \"record\", \"name\": \"Payment\", \"namespace\": \"com.tos\", \"fields\": [{\"name\": \"id\", \"type\": \"string\"}, {"name": \"amount\", \"type\": \"double\"}]}"} [root@kafka0 code]# █
```

#curl -X GET <http://localhost:8081/subjects>

Delete all schema versions registered under the subject “my-kafka-value”

curl -X DELETE http://localhost:8081/subjects/my-kafka-value

----- Lab Ends Here -----

9. Errors

I. {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)

```
[2018-05-15 23:46:40,132] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 14 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
[2018-05-15 23:46:40,266] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 15 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
^C[2018-05-15 23:46:40,394] WARN [Producer clientId=console-producer] Error whil
e fetching metadata with correlation id 16 : {test=LEADER_NOT_AVAILABLE} (org.apa
che.kafka.clients.NetworkClient)
[root@tos opt]# {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkCl
ient)
bash: syntax error near unexpected token `org.apache.kafka.clients.NetworkClient
'
```

Solutions: /opt/kafka/config/server.properties

Update the following information.

```
# it uses the value for "listeners" if configured. Otherwise, it will use the v
alue
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092
# Many listeners expose security protocols, the default is for them to be the s
ame
```

II Unable to connect to the server / Topic my-example-topic not present in
metadata after 60000 ms.

```
##### Socket Server Settings #####
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
#   FORMAT:
#     listeners = listener_name://host_name:port
#   EXAMPLE:
#     listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092

# Hostname and port the broker will advertise to producers and consumers. If not set,
# it uses the value for "listeners" if configured. Otherwise, it will use the value
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092

# Maps listener names to security protocols, the default is for them to be the same. See the config documentation for more details
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,SSL:SSL,SASL_PLAINTEXT:SASL_PLAINTEXT,SASL_SSL:SASL_SSL

"server.properties" 136L, 6848C written
```

10. LOG verification + segment Sizing

11. Annexure Code:

II. DumpLogSegment

```
/opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
/tmp/kafka-logs/my-kafka-connect-0/oooooooooooooooooooo.log | head -n 4
```

```
[root@tos test-topic-0]# more 00000000000000000000.log
[root@tos test-topic-0]# cd ..
[root@tos kafka-logs]# cd my-kafka-connect-0/
[root@tos my-kafka-connect-0]# ls
00000000000000000000.index      00000000000000000011.snapshot
00000000000000000000.log        leader-epoch-checkpoint
00000000000000000000.timeindex
[root@tos my-kafka-connect-0]# more *log
\kafka Connector.--More-- (53%)
```



```
[root@tos my-kafka-connect-0]# pwd
/tmp/kafka-logs/my-kafka-connect-0
[root@tos my-kafka-connect-0]# /opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
> /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log | head -n 4
Dumping /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log
Starting offset: 0
offset: 0 position: 0 CreateTime: 1530552634675 isvalid: true keysize: -1 valuesize: 31 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: This Message is from Test File .
offset: 1 position: 0 CreateTime: 1530552634677 isvalid: true keysize: -1 valuesize: 43 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: It will be consumed by the Kafka Connector.
[root@tos my-kafka-connect-0]#
```

CLI to list offset number.

```
#/opt/kafka/bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list localhost:9092  
--topic IBM
```

Sending Message with Key and Value

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic IBM --  
property "parse.key=true" --property "key.separator=:"
```

```
#/opt/kafka/bin/Kafka-console-consumer.sh --topic IBM --bootstrap-server  
localhost:9092 --from-beginning \  
--property print.key=true \  
--property key.separator=":" \  
--partition 4 \  
--offset 3
```

```
kafka-console-consumer --topic example-topic --bootstrap-server broker:9092 \  
--from-beginning \  
--property print.key=true \  
--property key.separator="-" \  
--partition 0
```

III. Resources

<https://developer.ibm.com/hadoop/2017/04/10/kafka-security-mechanism-saslplain/>

<https://sharebigdata.wordpress.com/2018/01/21/implementing-sasl-plain/>

<https://developer.ibm.com/code/howtos/kafka-authn-authz>

<https://jaceklaskowski.gitbooks.io/apache-kafka/content/kafka-tools-GetOffsetShell.html>